

Class 10: Structural Bioinformatics 1

Harshitha (PID: A17775151)

Table of contents

Background	1
PDB statistics	1
Visualizing PDB data with Mol-star	5
Getting to know HIV-Pr	6
Delving deeper	8
The important role of water	8
Getting started with the Bio3D package	10
Quick PDB visualization in R	12
Predict the flexibility of a given structure	12
Comparative analysis of the ADK family	16
Setup	16
Search and retrieve ADK structures	16
Align and superpose structures	22
Optional: Viewing our superposed structures	24
Annotate collected PDB structures	24
Principal component analysis	26
PCA visualization	28

Background

The main repository of high-resolution structural data on biomolecules is called the **Protein Data Bank** (PDB).

PDB statistics

What is in the PDB in terms of molecule type and structure determination method?

Read a CSV file of current PDB stats obtained from <https://www.rcsb.org/stats/summary>

```

pdb <- read.csv("Data Export Summary.csv")
pdb

```

	Molecular.Type	X.ray	EM	NMR	Integrative	Multiple.methods
1	Protein (only)	180,758	23,111	12,813	348	229
2	Protein/Oligosaccharide	10,488	3,741	34	8	11
3	Protein/NA	9,205	6,751	287	26	8
4	Nucleic acid (only)	3,154	250	1,578	3	15
5	Other	178	27	35	4	0
6	Oligosaccharide (only)	11	0	6	0	1
	Neutron	Other	Total			
1	84	32	217,375			
2	1	0	14,283			
3	0	0	16,277			
4	3	1	5,004			
5	0	0	244			
6	0	4	22			

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy.

80.49% of structures in the PDB are solved by X-ray, and 13.38% are solved by Electron Microscopy.

```

pdb$X.ray

```

```

[1] "180,758" "10,488"  "9,205"   "3,154"   "178"     "11"

```

This print out above `pdb$X.ray` is “character” not “numeric”. Therefore I can’t do math with it.

Two functions that can help here are `sub()` and `as.numeric()`

```

# We want to get rid (or sub out) commas:
x <- pdb$X.ray
tmp <- sub(",", "", x)
sum( as.numeric(tmp) )

```

```

[1] 203794

```

We could make a function to do this:

```
rm.comma <- function(x) {
  tmp <- sub(",", "", x)
  sum( as.numeric(tmp) )
}
```

```
# Percentage of structures solved by X-ray
rm.comma(pdb$X.ray) / rm.comma(pdb$Total)
```

```
[1] 0.8048577
```

```
# Percentage of structures solved by EM
rm.comma(pdb$EM) / rm.comma(pdb$Total)
```

```
[1] 0.1338046
```

We could also use a different import function for this CSV that speaks American (i.e. deals with commas in numbers in a comma separated value file).

```
library(readr)

pdb <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 9
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (1): Molecular Type
```

```
dbl (4): Integrative, Multiple methods, Neutron, Other
```

```
num (4): X-ray, EM, NMR, Total
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
pdb
```

```
# A tibble: 6 x 9
```

	`Molecular Type` <chr>	`X-ray` <dbl>	EM <dbl>	NMR <dbl>	Integrative <dbl>	`Multiple methods` <dbl>	Neutron <dbl>
1	Protein (only)	180758	23111	12813	348	229	84
2	Protein/Oligosacch~	10488	3741	34	8	11	1

3 Protein/NA	9205	6751	287	26	8	0
4 Nucleic acid (only)	3154	250	1578	3	15	3
5 Other	178	27	35	4	0	0
6 Oligosaccharide (o~	11	0	6	0	1	0

i 2 more variables: Other <dbl>, Total <dbl>

```
n.tot <- sum(pdb$Total)
n.xray <- sum(pdb$`X-ray`)
n.em <- sum(pdb$EM)

# Percentage of structures solved by X-ray
(n.xray / n.tot) * 100
```

[1] 80.48577

```
# Percentage of structures solved by EM
(n.em / n.tot) * 100
```

[1] 13.38046

Q. How many total protein structures are there in the dataset?

There are 217375 protein structures in the database.

```
pdb$Total[1]
```

[1] 217375

Q2: What proportion of structures in the PDB are protein?

0.107% of PDB structures are protein.

```
# The total number of entries in UniProt is 202,556,314.
(pdb$Total[1] / 202556314) * 100
```

[1] 0.1073158

Key-point: We have a very, very small structural coverage of known proteins (~0.1%). Most structures we know about (~80%) come from one method (X-ray crystallography).

SKIPPED IN CLASS Q3: Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB?

Visualizing PDB data with Mol-star

Main stand alone web version with all features is at <https://molstar.org/viewer/>



Figure 1: The HIV-protease enzyme is a homodimer with two chains

Getting to know HIV-Pr



Figure 2: Ligand with Spacefill representation, and polymer colored by secondary structure




Figure 3: Surface display showing the binding cleft site for the inhibitor (drug)

Delving deeper



Figure 4: Display showing the Asp 25 (D25) residue in each chain, which is critical to protease activity. The polymer is colored by secondary structure.

The important role of water

Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?

We only see each water molecule's oxygen atom. X-ray crystallography has a 2.00 Angstrom resolution, and hydrogen is smaller than this – thus, it does not appear in the structure.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have?

The water molecule has the residue number 308.

Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic

residues ASP 25 in each chain and the critical water. Add this figure to your Quarto document.

Discussion Topic: Can you think of a way in which indinavir, or even larger ligands and substrates, could enter the binding site?



Figure 5: Spacefill display of catalytic ASP25 amino acids and key binding site water molecule. Ligand is displayed with a ball-and-stick representation.

For indinavir, or other ligands/substrates to enter the binding site, it's possible that the two loops at the top of the protease structure (as shown in the image above; one loop per chain) undergo a conformational change and “open” in a gate-like manner to enable substrate entry. Once the substrate enters, the loops return to their original, “closed” conformation. The positioning of the loops and the newly-entered substrate may be stabilized by non-covalent/hydrogen bonding interactions with the critical water molecule (HOH 308, shown in the image above). Furthermore, the substrate is stabilized by interactions with the catalytic Asp 25 residues in the protease's active site.

Getting started with the Bio3D package

The following code has been inputted into the R brain, to install useful packages. Pak can install from multiple sources (i.e. “bioboot/bio3dview” from GitHub, “NGLVieweR” from CRAN, “bioc::msa” from Bioconductor).

```
# install.packages("pak")
# pak::pkg_install( c("bioboot/bio3dview", "NGLVieweR", "bioc::msa"))
```

Bio3D is an R package from CRAN for structural bioinformatics.

```
library(bio3d)

pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
pdb
```

```
Call: read.pdb(file = "1hsg")
```

```
Total Models#: 1
  Total Atoms#: 1686,  XYZs#: 5058  Chains#: 2  (values: A B)

Protein Atoms#: 1514  (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0  (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172  (residues: 128)
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

```
Protein sequence:
```

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,
      calpha, remark, call
```

Q7: How many amino acid residues are there in this pdb object?

There are 198 amino acid residues.

Q8: Name one of the two non-protein residues?

One non-protein residue is HOH (res no. 127). The other is MK1 (res no. 1).

Q9: How many protein chains are in this structure?

There are two protein chains.

```
# To access dataset's attributes
```

```
attributes(pdb)
```

```
$names
```

```
[1] "atom" "xyz" "seqres" "helix" "sheet" "calpha" "remark" "call"
```

```
$class
```

```
[1] "pdb" "sse"
```

```
# Access the atom attribute in pdb
```

```
head(pdb$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40

	segid	elesy	charge
1	<NA>	N	<NA>
2	<NA>	C	<NA>
3	<NA>	C	<NA>
4	<NA>	O	<NA>
5	<NA>	C	<NA>
6	<NA>	C	<NA>

There are lots of functions that can work with these `pdb` objects:

```
head(pdbseq(pdb))
```

```
1 2 3 4 5 6  
"P" "Q" "I" "T" "L" "W"
```

Quick PDB visualization in R

We can have a quick view of any of these `pdb` objects:

```
library(bio3dview)
library(NGLViewerR)

view.pdb(pdb) |>
  setSpin()
```

Let's try a custom view:

```
# To color by secondary structure
view.pdb(pdb,
  colorScheme = "sse",
  backgroundColor = "black")
```

Q. Create a custom view of HIV-Pr highlighting the active site ASP residues (`resno=25`), the two chains (in your choice of colors), and the ligand all on a custom color background.

```
# Select the ASP 25 residue
active.site <- atom.select(pdb, resno = 25)

# Highlight D25 residues in spacefill representation
view.pdb(pdb,
  cols = c("green", "teal"),
  highlight = active.site,
  highlight.style = "spacefill",
  backgroundColor = "black") |>
  setRock()
```

Predict the flexibility of a given structure

Let's do a Normal Mode Analysis (NMA) to predict the flexibility of a given `pdb` object:

```
# Step 1: Read in PDB structure of Adenylate Kinase
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file

PDB has ALT records, taking A only, `rm.alt=TRUE`

```
adk
```

```
Call: read.pdb(file = "6s36")
```

```
Total Models#: 1
```

```
Total Atoms#: 1898, XYZs#: 5694 Chains#: 1 (values: A)
```

```
Protein Atoms#: 1654 (residues/Calpha atoms#: 214)
```

```
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 244 (residues: 244)
```

```
Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]
```

```
Protein sequence:
```

```
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV  
DELVIALVKERIAQEDCRNGFLLDGFPRTPQADAMKEAGINVDYVLEFDVPDELIVDKI  
VGRRVHAPSGRVYHVKFNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTAPLIG  
YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
```


```
+ attr: atom, xyz, seqres, helix, sheet,  
      calpha, remark, call
```

```
# Step 2: Pass pdb object to nma() function  
m <- nma(adk)
```

```
Building Hessian... Done in 0.017 seconds.
```

```
Diagonalizing Hessian... Done in 0.413 seconds.
```

```
# Step 3a: Pass nma results to plot(). Fluctuations plot is particularly interesting!  
plot(m)
```



View the results with an interactive structure view (use `view.nma()`):

```
# Step 3b: Pass nma results to view.nma()
view.nma(m)
```

Write out the results for viewing in Mol-star:

```
# Step 3c: Pass nma results to mktrj(), to generate file to open in Mol-star
mktrj(m, file="nma.pdb")
```



Figure 6: Mol-star view of Adenylate Kinase. Playing the movie in Mol* shows a highly dynamic, “opening/closing” motion of the structure.

Comparative analysis of the ADK family

Setup

The following installations are relevant for this lab, and were already done in the R console:

```
# install.packages("bio3d")
# install.packages("NGLVieweR")

# install.packages("remotes")
# remotes::install_github("bioboot/bio3dview")

# install.packages("BiocManager")
# BiocManager::install("msa")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

“msa” is found only on BioConductor (must be installed using `BiocManager::install()`).

Q11. Which of the above packages is not found on BioConductor or CRAN?:

“bio3dview” is not found on BioConductor or CRAN (found on GitHub instead, and must be installed using `devtools::install_github()`).

Q12. True or False? Functions from the pak package can be used to install packages from GitHub and BitBucket?

True! Pak can install from multiple sources.

Search and retrieve ADK structures

Our first step is find a sequence for this family. We will use the database ID “1ake_A” here:

```
id <- "1ake_A"

# Step 1: Use get.seq() with input database identifier.
aa <- get.seq(id)
```

Warning in `get.seq(id)`: Removing existing file: `seqs.fasta`

Fetching... Please wait. Done.


```
aa
```

```
      1      .      .      .      .      .      .      60
pdb|1AKE|A  MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV
      1      .      .      .      .      .      .      60

      61      .      .      .      .      .      .      120
pdb|1AKE|A  DELVIALVKERIAQEDCRNGFLLDGFPRTPQADAMKEAGINVDYVLEFDVPDELIVDRI
      61      .      .      .      .      .      .      120

      121      .      .      .      .      .      .      180
pdb|1AKE|A  VGRRVHAPSGRVYHVKNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTPALIG
      121      .      .      .      .      .      .      180

      181      .      .      .      214
pdb|1AKE|A  YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
      181      .      .      .      214
```

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

```
+ attr: id, ali, call
```

Q13. How many amino acids are in this sequence, i.e. how long is this sequence?

There are 214 amino acids in the sequence.

Search for related sequences in the database

```
# Step 2a: Find related sequences via a PDB search, using blast.pdb() function
blast <- blast.pdb(aa)
```

```
Searching ... please wait (updates every 5 seconds) RID = 020ZJ0B9014
```

```
...
```

```
Reporting 96 hits
```

```
# Preview the top hits
head(blast$hit.tbl)
```


	queryid	subjectids	identity	alignmentlength	mismatches	gapopens	q.start
1	Query_229329	1AKE_A	100.000	214	0	0	1
2	Query_229329	8BQF_A	99.533	214	1	0	1
3	Query_229329	4X8M_A	99.533	214	1	0	1
4	Query_229329	6S36_A	99.533	214	1	0	1
5	Query_229329	9R6U_A	99.533	214	1	0	1
6	Query_229329	9R71_A	99.533	214	1	0	1

	q.end	s.start	s.end	evaluate	bitscore	positives	mlog.evaluate	pdb.id	acc
1	214	1	214	1.82e-156	432	100.00	358.6044	1AKE_A	1AKE_A
2	214	21	234	2.97e-156	433	100.00	358.1147	8BQF_A	8BQF_A
3	214	1	214	3.25e-156	432	100.00	358.0246	4X8M_A	4X8M_A
4	214	1	214	4.77e-156	432	100.00	357.6409	6S36_A	6S36_A
5	214	1	214	1.06e-155	431	99.53	356.8424	9R6U_A	9R6U_A
6	214	1	214	1.25e-155	431	99.53	356.6775	9R71_A	9R71_A

```
# Plot BLAST results
hits <- plot(blast)
```

```
* Possible cutoff values:    260 3
    Yielding Nhits:         20 96

* Chosen cutoff value of:    260
    Yielding Nhits:         20
```



```
# Each dot is one row from hits table.
# Dashed line represents point of largest drop-off in norm. scores (E-value jump)
# Black dots to the left of dashed line are considered good hits
# Red dots to the right of dashed line have high E-values
```

```
# List out the top hits
hits$ pdb.id
```

```
[1] "1AKE_A" "8BQF_A" "4X8M_A" "6S36_A" "9R6U_A" "9R71_A" "8Q2B_A" "8RJ9_A"
[9] "6RZE_A" "4X8H_A" "3HPR_A" "1E4V_A" "5EJE_A" "1E4Y_A" "3X2S_A" "6HAP_A"
[17] "6HAM_A" "8PVW_A" "4K46_A" "4NP6_A"
```

```
#Step 2b: Download BLAST hits with get.pdb() function.
files <- get.pdb(hits$ pdb.id, path="pdbs", split=TRUE, gzip=TRUE)
```

```
Warning in get.pdb(hits$ pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/1AKE.pdb.gz exists. Skipping download
```

```
Warning in get.pdb(hits$ pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/8BQF.pdb.gz exists. Skipping download
```

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/4X8M.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6S36.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/9R6U.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/9R71.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/8Q2B.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/8RJ9.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6RZE.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/4X8H.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/3HPR.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/1E4V.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/5EJE.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/1E4Y.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/3X2S.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6HAP.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/6HAM.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/8PVW.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/4K46.pdb.gz exists. Skipping download

Warning in get.pdb(hits\$pdb.id, path = "pdbs", split = TRUE, gzip = TRUE):
pdbs/4NP6.pdb.gz exists. Skipping download

	0%
====	5%
=====	10%
=====	15%
=====	20%
=====	25%
=====	30%
=====	35%
=====	40%
=====	45%
=====	50%
=====	55%

=====	60%
=====	65%
=====	70%
=====	75%
=====	80%
=====	85%
=====	90%
=====	95%
=====	100%

Align and superpose structures

Align and superpose all these ADK structures

```
# Step 3: Align related PDBs using pdbaln() function, in superposed manner
pdbs <- pdbaln(files, fit = TRUE, exefile="msa")
```

Reading PDB files:

```
pdbs/split_chain/1AKE_A.pdb
pdbs/split_chain/8BQF_A.pdb
pdbs/split_chain/4X8M_A.pdb
pdbs/split_chain/6S36_A.pdb
pdbs/split_chain/9R6U_A.pdb
pdbs/split_chain/9R71_A.pdb
pdbs/split_chain/8Q2B_A.pdb
pdbs/split_chain/8RJ9_A.pdb
pdbs/split_chain/6RZE_A.pdb
pdbs/split_chain/4X8H_A.pdb
pdbs/split_chain/3HPR_A.pdb
pdbs/split_chain/1E4V_A.pdb
pdbs/split_chain/5EJE_A.pdb
pdbs/split_chain/1E4Y_A.pdb
pdbs/split_chain/3X2S_A.pdb
```

```

pdbs/split_chain/6HAP_A.pdb
pdbs/split_chain/6HAM_A.pdb
pdbs/split_chain/8PVW_A.pdb
pdbs/split_chain/4K46_A.pdb
pdbs/split_chain/4NP6_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
..   PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
..   PDB has ALT records, taking A only, rm.alt=TRUE
..   PDB has ALT records, taking A only, rm.alt=TRUE
....  PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
.    PDB has ALT records, taking A only, rm.alt=TRUE
..

```

Extracting sequences

```

pdb/seq: 1    name: pdbs/split_chain/1AKE_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 2    name: pdbs/split_chain/8BQF_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 3    name: pdbs/split_chain/4X8M_A.pdb
pdb/seq: 4    name: pdbs/split_chain/6S36_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 5    name: pdbs/split_chain/9R6U_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 6    name: pdbs/split_chain/9R71_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 7    name: pdbs/split_chain/8Q2B_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 8    name: pdbs/split_chain/8RJ9_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 9    name: pdbs/split_chain/6RZE_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 10   name: pdbs/split_chain/4X8H_A.pdb
pdb/seq: 11   name: pdbs/split_chain/3HPR_A.pdb
    PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 12   name: pdbs/split_chain/1E4V_A.pdb

```

```

pdb/seq: 13  name: pdbs/split_chain/5EJE_A.pdb
  PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 14  name: pdbs/split_chain/1E4Y_A.pdb
pdb/seq: 15  name: pdbs/split_chain/3X2S_A.pdb
pdb/seq: 16  name: pdbs/split_chain/6HAP_A.pdb
pdb/seq: 17  name: pdbs/split_chain/6HAM_A.pdb
  PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 18  name: pdbs/split_chain/8PVW_A.pdb
  PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 19  name: pdbs/split_chain/4K46_A.pdb
  PDB has ALT records, taking A only, rm.alt=TRUE
pdb/seq: 20  name: pdbs/split_chain/4NP6_A.pdb

```

```
# pdbs
```

Optional: Viewing our superposed structures

Quick interactive structural view

```
view.pdbs(pdbs)
```

Annotate collected PDB structures

We can use `pdb.annotate()` to annotate PDB files.

```

# Vector with PDB database codes of collected structures
ids <- basename.pdb(pdbs$id)

# Here, we use pdb.annotate() to annotate structure to source species.
anno <- pdb.annotate(ids)
unique(anno$source)

```

```

[1] "Escherichia coli"
[2] "Escherichia coli K-12"
[3] "Escherichia coli 0139:H28 str. E24377A"
[4] "Escherichia coli str. K-12 substr. MDS42"
[5] "Photobacterium profundum"
[6] "Vibrio cholerae 01 biovar El Tor str. N16961"

```



```
# Columns of anno are the available annotation features.
head(anno)
```

	structureId	chainId	macromoleculeType	chainLength	experimentalTechnique
1AKE_A	1AKE	A	Protein	214	X-
ray					
8BQF_A	8BQF	A	Protein	234	X-
ray					
4X8M_A	4X8M	A	Protein	214	X-
ray					
6S36_A	6S36	A	Protein	214	X-
ray					
9R6U_A	9R6U	A	Protein	214	X-
ray					
9R71_A	9R71	A	Protein	214	X-
ray					

	resolution	scopDomain	pfam	ligandId
1AKE_A	2.00	Adenylate kinase	Adenylate kinase (ADK)	AP5
8BQF_A	2.05	<NA>	Adenylate kinase (ADK)	AP5
4X8M_A	2.60	<NA>	Adenylate kinase (ADK)	<NA>
6S36_A	1.60	<NA>	Adenylate kinase (ADK)	CL (3),NA,MG (2)
9R6U_A	1.77	<NA>	Adenylate kinase (ADK)	AP5,GOL,NA
9R71_A	1.61	<NA>	Adenylate kinase (ADK)	AP5

	ligandName	source
1AKE_A	BIS(ADENOSINE)-5'-PENTAPHOSPHATE	Escherichia coli
8BQF_A	BIS(ADENOSINE)-5'-PENTAPHOSPHATE	Escherichia coli
4X8M_A	<NA>	Escherichia coli
6S36_A	CHLORIDE ION (3),SODIUM ION,MAGNESIUM ION (2)	Escherichia coli
9R6U_A	BIS(ADENOSINE)-5'-PENTAPHOSPHATE,GLYCEROL,SODIUM ION	Escherichia coli
9R71_A	BIS(ADENOSINE)-5'-PENTAPHOSPHATE	Escherichia coli

1AKE_A STRUCTURE OF THE COMPLEX BETWEEN ADENYLATE KINASE FROM ESCHERICHIA COLI AND THE INHIBITORS
8BQF_A
4X8M_A
6S36_A
9R6U_A
9R71_A

	citation	rObserved	rFree
1AKE_A	Muller, C.W., et al. J Mol Biology (1992)	0.19600	NA
8BQF_A	Scheerer, D., et al. Proc Natl Acad Sci U S A (2023)	0.22073	0.25789
4X8M_A	Kovermann, M., et al. Nat Commun (2015)	0.24910	0.30890
6S36_A	Rogne, P., et al. Biochemistry (2019)	0.16320	0.23560


9R6U_A		Mattsson, J., et al. Biochemistry (2025)	NA	0.22790
9R71_A		Mattsson, J., et al. Biochemistry (2025)	0.19600	0.24400
	rWork	spaceGroup		
1AKE_A	0.19600	P 21 2 21		
8BQF_A	0.21882	P 2 21 21		
4X8M_A	0.24630	C 1 2 1		
6S36_A	0.15940	C 1 2 1		
9R6U_A	0.19190	P 21 2 21		
9R71_A	0.19300	P 21 21 2		

Principal component analysis


PCA of all this structural data (x, y, and z atom coordinates):

```
# Step 4: Run PCA analysis for aligned structural data
pc <- pca(pdfs)

# Generate score plots and loadings plots
plot(pc)
```



```
# Generate score plot of PC1 vs PC2
plot(pc, 1:2)
```




We can cluster structures using RMSD values to determine pairwise structural deviations:

```
# Calculate the RMSD values
rd <- rmsd(pdb)
```

Warning in rmsd(pdb): No indices provided, using the 182 non NA positions

```
# Group structures into 3 clusters using RMSD values
hc.rd <- hclust(dist(rd))
grps.rd <- cutree(hc.rd, k=3)

# Plot PC1 vs PC2, w/ points colored by RMSD-based clustering assignment
plot(pc, 1:2, col="grey50", bg=grps.rd, pch=21, cex=1)
```



The plot above depicts the separation and grouping of our collected PDB structures along PC1 and PC2, representing conformational differences among the structures.

PCA visualization

Interactive view of the PCA-captured structural differences:

```
view.pca(pc)
```

To view structural variation along PC1, in Mol*, first use `mktrj()` to generate a trajectory PDB file. The file can then be opened in Mol-star:

```
mktrj(pc, pc = 1, file = "pc_1.pdb")
```



Figure 7: Image of Mol-star animation of structural variation along PC1.